



Dream Job

Functional walkthrough and technical architecture · v2

An AI-assisted job discovery and application platform: profile in, ranked and explained opportunities out, every message reviewed and sent by the job seeker. Version 2 refreshes every screenshot, diagram and count, and covers the screens added since the first edition.

Prepared for Stephane van der Aa | Live application walkthrough, system architecture and the full source catalogue · September 2026

What this deck contains

A working tour of the platform, the architecture behind it, and the complete list of external sources it can read.

- Part 1 · A functional walkthrough: a screenshot of every working screen — including Browser session, Apply browser, Networking and Mail setup, which the first edition did not cover — with what each screen is for and how the job seeker uses it.
- Part 2 · Technical architecture: the layered system, the five chokepoints, the campaign pipeline end to end, the adapter contract, the path of one LLM call and one outbound request, and the data model.
- Part 3 · Data sources: every registered adapter — 29 today — grouped by type, with coverage, access method, terms status and whether it needs the language model.
- Throughout: diagrams and figures are generated from the code, and every count in this deck was measured from the working tree.
- The principle throughout: nothing is sent or applied on its own; the job seeker is the last line of defence and the only sender.

What is new in version 2


Every figure below was re-measured from the current codebase; screens added since the first edition are listed on the right.

Metric	v1	v2 (measured)
Python modules	172	227 (120,881 lines)
JS/JSX modules	113	140 (43,932 lines)
API routes	301	422 over 22 routers
Database tables	71	88 (49 private, 38 shared)
Indexes	144	158
Migrations	34	43
Tests	1,725	2,055 passing of 2,061 (114 files)
Adapters	29	29 (4 need acknowledgement)
Prompt templates	24	24

- New screens: Browser session, Apply browser, Networking and export, Mail setup, and the per-screen Help drawer.
- New depth within screens: Profile tabs (Sections, Skills, Evidence, Personas, Privacy), Composite (Online findings, Enrichment), Campaign (Plan review, Live dashboard, Re-run a stage), Company (Profile, Financials, Market and timing, Values and culture).
- New depth continued: Contacts (Browse all, Introduction routes, My network, Objections and retention), Dream-job intelligence (Gaps, Stepping stones, Fit across the market, Values conflicts, LinkedIn profile), Administration (nine tabs).
- Refreshed diagrams: all 20 generated figures and every screenshot in this deck come from the current system.

The product in one picture

- Profile — import the facts about you and describe the job you want.
- Plan — bound the search with structured directives and review the plan.
- Discover — collect, profile, analyse and rank the market.
- Apply — write, check and dispatch applications you approve.
- Follow up — read the outcomes and redirect the next search.



cool: working out what you want

warm: where messages leave the building

A phase keeps its hue on every screen it owns. The spectrum runs cool to warm because the stakes do: the first phases only decide what you want, and nothing reaches a stranger until the warm end.

The five phases of the journey. Each phase owns a colour that follows it through every screen of the application.

PART 1


A functional walkthrough

Every screen, in the order a job seeker meets it. The screenshots are taken from the running application with real data.

Start here

Three steps and nothing else: set up your profile, let Dream Job search and rank the market, then choose who to write to. Everything in between runs on its own.

- The whole search runs in the background; the page can be closed.
- Live counts on today's data: 1,084 opportunities found, 21 speculative, 939 ranked, 4 with a contact.
- Nothing is ever sent on its own — every message waits for approval.
- A speculative opening is labelled everywhere so an email never claims a vacancy exists.




home · three steps, one action each

Where I am — the full journey map

The whole process on one screen: what is done, what is running, what is blocked, and what the next useful move is. Every stage links to the screen that advances

- A green mark is finished, a pulsing one is running, a dotted one is waiting on an earlier stage.
- Headline counts are links to the screens their numbers come from.
- While a campaign runs, its token budget and cost refresh every few seconds.
- A blocked stage says what it needs, so the next move is never a guess.




overview · the journey map over the live campaign

Your profile

Everything the system does starts here. The profile is the only source of facts about the job seeker; nothing in a generated CV or email can say anything not


- Import the LinkedIn export and the CV; the two are merged, and conflicts are surfaced, not silently resolved (FR-102/103).
- This account holds 22 saved versions; every save is a new version and campaigns record which version they used.
- Five extra tabs sit one click away: Sections, Skills, Evidence, Personas and Privacy.
- Fields marked 'do not disclose' are stripped before anything reaches the AI provider (FR-106).




profile · Documents — the two uploads

Resolving the disagreements


The merge is honest about what it could not decide: every conflict is shown with both values side by side and stays open until the job seeker chooses.




profile · Conflicts — contested values, resolved by hand



profile · Skills — the normalised, canonical skill taxonomy




profile · Evidence — statements kept with their provenance




profile · Privacy — fields marked do not disclose

The rest of the profile workspace


Personas let one profile serve more than one kind of application; sections show the structured profile the model actually reads.




profile · Sections — the structured profile



profile · Personas — one profile, more than one angle



profile · Dream job tab — a summary with a link to the editor




profile · the default view, with version and conflict counters

Composite profile

A synthesis of what the job seeker supplied and what could be verified online. This is what the model actually reads when matching, scoring and writing.


- Every statement carries its source and a confidence.
- Ten blocks from narrative and career trajectory to constraints and inferred preferences.
- Online findings are queued for confirm or reject; a rejection is permanent (FR-124).
- Enrichment can be switched off; the profile is then built from documents alone.




composite · the traceable synthesis

Online findings and enrichment


Identity matching folds name variants, employer overlap, cross-links, location, timeline and photo similarity; only confirmed findings merge (RK-02).




composite · Online findings — confirm or reject, permanently



composite · Enrichment — switch it off, exclude domains



composite · the composite in context



the same discipline as profile conflicts

Dream job

Describe the job you actually want, in your own words and language. The model reads it for nuance — target roles, responsibilities, values and deal-breakers.

- Deal-breakers filter more decisively than preferences; they become directive defaults.
- The structured model is used for matching; the prose is not.
- The screen flags when the statement and the model have drifted apart.
- Nothing acts on the model until the job seeker confirms it (FR-128).


The screenshot shows the 'Dream job' interface. On the left is a sidebar with navigation links: 'DO THIS' (Start here, 1 - Your profile, 2 - Opportunities, 3 - Apply, Follow up), 'ADVANCED' (Composite profile, **Dream job**, Directives, Campaigns, Browser session), 'RESEARCH & APPLY' (Companies, Dream-job gap, Contacts, Apply browser, Networking), 'JOURNEY & FOLLOW UP' (Where I am, Responses, What works), and 'SYSTEM' (Monitoring, Mail setup, Administration). The main content area is titled 'Dream job' and includes a header with 'Companies 5,941', 'Jobs 57,293', 'Contacts 3,959', and user info 'stephane@stepvda.com'. Below the header is a progress bar with 'Profile' (active), 'Plan', 'Discover', 'Apply', and 'Follow up'. The 'Your statement' section contains a text area with a 133-word limit and a 'Rebuild the model from this text' button. Below this is the 'Structured model, version 11' section, which is marked 'Not confirmed' and dated '13 Sept 2026'. It includes a confirmation button and a note about the extraction process. The structured model is divided into four panels: 'Target roles' (Software Architect, priority 1, stated), 'Role families' (hands-on senior software engineering, inferred from your statement), 'Responsibilities' (Writing production code, must, stated), and 'Company characteristics' (Brussels or remote; no relocation, must, stated).

dream job · statement plus structured model

Search directives

Directives bound the search. They are structured controls rather than free text, so they translate reliably into each source's own query language.

- Five groups: job content, company type, location, work arrangement and compensation.
- Values pre-filled from the composite profile are suggestions, not facts.
- Discretion mode excludes the current employer and its group entities from every search and message (FR-385).
- The estimate beside the form says how many sources and pages the settings imply.



directives · structured controls and a live estimate

Campaigns

A campaign is one run of the collection and analysis pipeline under one set of directives. Every stage stores what it produced and can be re-run on its own.

- The plan is reviewed before launch: sources, queries, expected volume, duration and cost.
- The planner checks the shared knowledge base first and collects only what is missing or stale (FR-342/343).
- Pause, resume or cancel at any point; a crash loses at most the page in flight (NFR-401).
- The token budget caps AI spend and degrades gracefully near the limit.

Dream Job
Stephane van der Aa

DO THIS

- Start here
- 1 · Your profile
- 2 · Opportunities
- 3 · Apply
- Follow up

ADVANCED

PROFILE & SEARCH

- Composite profile
- Dream job
- Directives
- Campaigns**
- Browser session

RESEARCH & APPLY

- Companies
- Dream-job gap
- Contacts
- Apply browser
- Networking

JOURNEY & FOLLOW UP

- Where I am
- Responses
- What works

SYSTEM

- Monitoring
- Mail setup
- Administration

Campaigns

Companies 5,941 Jobs 57,293 Contacts 3,959 stephane@stepvda.com ? Help Sign out

Profile Plan **Discover** Apply Follow up

A campaign is one run of the collection and analysis pipeline under one set of directives. Each stage stores what it produced, so any stage can be re-run on its own.

+ New campaign

CAMPAIGN	STATUS	% DIRECTIVE SET	% STAGE	% TOKENS USED	\$ COST	TIMING	
Continuous data collection created 13 Sept 2026	completed	Autopilot directives	collection	891,335 / 2,000,000	€0.36	13 Sept 2026 ran for 11 min	Delete
Search 13 Sept created 13 Sept 2026	completed	Autopilot directives	collection	740,990 / 2,000,000	€0.31	13 Sept 2026 ran for 55 min	Delete
Autopilot 12 Sep created 13 Sept 2026	completed	Autopilot directives	collection	769,486 / 2,000,000	€0.31	13 Sept 2026 ran for 12 min	Delete
Autopilot 12 Sep created 12 Sept 2026	completed	Autopilot directives	collection	415,360 / 2,000,000	€0.16	12 Sept 2026 ran for 9 min	Delete
Autopilot 12 Sep created 12 Sept 2026	completed	Autopilot directives	collection	1,210,056 / 2,000,000	€0.47	12 Sept 2026 ran for 11 min	Delete
Autopilot 12 Sep created 12 Sept 2026	completed	Autopilot directives	collection	17,858 / 2,000,000	€0.01	12 Sept 2026 ran for 7 min	Delete
Search 12 Sept created 12 Sept 2026	completed	Rebuilt from profile 2026-09-11	collection	1,995,075 / 2,000,000	€0.78	12 Sept 2026 ran for 21h 37m	Delete
Search 11 Sept 8 created 11 Sept 2026	completed	Rebuilt from profile 2026-09-11	collection	1,713,626 / 2,000,000	€0.66	11 Sept 2026 ran for 41 min	Delete
Search 11 Sept created 11 Sept 2026	completed	Proposed directives	collection	393,675 / 2,000,000	€0.13	11 Sept 2026 ran for 2h 17m	Delete
Search 11 Sept created 11 Sept 2026	completed	Autopilot directives	collection	158,558 / 2,000,000	€0.06	12 Sept 2026 ran for 56 min	Delete
Autopilot 11 Sep created 11 Sept 2026	completed	Autopilot directives	collection	62,775 / 2,000,000	€0.02	11 Sept 2026 ran for 35s	Delete
Autopilot 11 Sep created 11 Sept 2026	completed	Autopilot directives	collection	508,392 / 2,000,000	€0.22	11 Sept 2026 ran for 2 min	Delete
Autopilot 11 Sep created 11 Sept 2026	completed	Autopilot directives	collection	41,615 / 2,000,000	€0.01	11 Sept 2026 ran for 48s	Delete
Search 11 Sept created 11 Sept 2026	completed	Proposed directives	collection	25,146 / 2,000,000	€0.01	11 Sept 2026 ran for 33 min	Delete

campaigns · the ledger of runs

Brain jobs
member since 10 yr

My tools

- Start here
- 1. Your profile
- 2. Openworkables
- 3. History
- Follow up

INTERESTS

PROFESSOR & RESEARCHER

- Entrepreneurial profile
- Brain job
- Discussions
- Disciplines**
- Research session

RESEARCH & ANALYSIS

- Comparative
- Discipline map
- Concepts
- Apply browser
- Networking

ANALYSIS & VISUALIZATION

- Where I am
- Regression
- What works
- Network
- Monitoring
- Mail setup
- Administration

Disciplines

At a glance **Brain** **Discipline** **History** **Follow-up**

• All disciplines • **Continuance data collection** **Completed** 1 collection

A discipline is one of the collections that are used to gather open workable disciplines. Each discipline stores what is produced, so any stage can be run or not on its own.

Filter modes: **list** **table** **discipline dashboard** **Discipline as a stage** **X**

Estimated for this plan stage: collection

106 disciplines to run

106 disciplines

14 min estimated

€ 0.00 estimated per discipline

Estimated by plan: 106 disciplines already have 106 disciplines to run, 106 disciplines to run, 14 min estimated, 0 estimated per discipline. **Nothing has been taken, find the parties, include what you do not want. Run again.**

Reused from the knowledge base Re-check

Knowledge base: 125 categories, internal model, 1000 to collect, 0 categorization, observation, internal model, 400 not to collect, 400 event, internal model, 0 not to collect, 1000 internal model, 100 with collect, 1000 not to collect, 1000 not to collect and 1000 not to collect.

11 min estimated per event

€ 0.00

7 days Run from 10/10/2016 to 16/10/2016 10:10

RECORD TYPE	FROM THE KNOWLEDGE BASE	REUSED FOR THIS CONCEPT	STILL TO COLLECT	RE-DELETED AFTER
category	2,000	300	0	30 days
categorization, observation	0	0	0	5400 days
event	0	400	0	30 days
category	10,100	5,000	100	7 days

0 SOURCE **DISCIPLINE** **REUSED** & **EXCLUDED** **PAGES** **1 DAY**

used events **0/1/0** **1 + 0** 00 have internal model from this source are within the 7-day retention period

used evidence **100/100** **1 + 0** 00 have internal model from this source are within the 7-day retention period

[illegible]

Stream job

Platform version 2.0

Start here

1. Your profile

2. Organization

3. Agency

Follow up

COMPONENTS

PROVIDER & SERVICES

Generic profile

Stream job

Streamlines

Components

Streamline job

Generic

Apply Streamer

Streamlining

Where can I find what works?

Monitoring

Mail setup

Administration

Components

Active

Plan

Streamers

Ready

Follow up

At campaigns

Continuous data collection

Import collection

Campaigns

Use one of the collection and analysis algorithms under one set of directives.

Each step starts with a condition, so the steps can flow from one to the next.

Plan review

100

Use dashboard

Monitor at stage

1

Be run one stage

1

Every stage starts with a condition, so a stage can run again and its set – register after adding your directives, to make after changing your weights – collect collecting anything else.

STAGE	What it does	
collect_plan	Re-run collection from the plan (20-30, 100)	Be run
company_profile	Substitutes the company profile from what is stored (20-30, 100)	Be run
learn_labels	Re-learn the things and the system ability to say (20-30, 100)	Be run
support_labels	Re-learn the support from the collected (20-30, 100)	Be run
learning	Re-learn the system plan (20-30, 100)	Be run
reset	Re-learn the knowledge from the plan (20-30, 100)	Be run
start	Re-learn the system capability in the campaign (20-30, 100)	Be run
specific_run	Re-learn the system ability for the campaign with the matching capacity (20-30, 100)	Be run

Components

6,501

1 job

\$7,335

3 Comments

streamlinetools.com

1 help

Sign out


[illegible]

Dream Job · Functional walkthrough and technical architecture · v2

Browser session

The control surface for browser-assisted collection. Dream Job attaches to a browser you already signed in to; it never holds a third-party password (FR-201/202).

- Connection panel reports whether a browser is attached and which sites are signed in.
- Launch instructions per operating system; the session is attach-only.
- Human pacing, a duration contract and immediate stop on a challenge (CR-401).
- Retention and purge keep browser-collected data on a short leash (NFR-303/344).



browser · connection, pacing and terms acknowledgement

Opportunities

Every advertised vacancy and every speculative opening the system found, ranked. The score is advisory; the job seeker decides what to pursue.

- Filter and sort by kind, arrangement, seniority, country, score floor or keyword.
- 'Why this rank?' opens the seven sub-scores and the dream-job criteria behind them.
- Manual ordering overrides the computed ranking and survives recalculation (FR-284).
- Rejection reasons feed back into the weights, which is why a reason is required.

The screenshot displays the 'Opportunities' section of the Dream Job application. The interface includes a sidebar with navigation options like 'DO THIS', 'PROFILE & SEARCH', 'RESEARCH & APPLY', 'JOURNEY & FOLLOW UP', and 'SYSTEM'. The main content area features a top navigation bar with 'Profile', 'Plan', 'Discover' (active), 'Apply', and 'Follow up'. Below this, a message states: 'Every advertised vacancy and every speculative opening the system found, ranked. The score is advisory; you decide what to pursue.'

A yellow box highlights: 'Scores are advisory. Scores are advisory. They order a list; they do not decide anything. Every decision to keep, reject or apply is the job seeker's (NFR-305, CR-405).'

The 'Campaign' section shows 'Continuous data collection · comple' and 'Sort' options: 'Overall score, best first' (selected), 'My order first', and a 'Find more opportunities' button. Below this, statistics are shown: 22 Opportunities, 22 Advertised vacancies, 0 Speculative openings, and 0 On my shortlist.

The 'Company or keyword' section includes a search bar with 'e.g. Colruyt, platform, Ghent' and filters for 'Kind' (Vacancies and speculative c), 'Tag' (Any tag), 'Timing' (Any timing), and 'Minimum score: 0'. A note states: 'Keyword matches the title, the description and the company name. There is no upper score bound — sort descending to see the strongest matches first.'

The 'Work arrangement' section has filters for 'Any work arrangement', 'Any contract', 'Any seniority', 'Any function', and 'Any country'. A checkbox for 'Hide the ones I rejected' is also present.

A blue box contains the instruction: 'Drag the # handle to place a row by hand. A row you place keeps its position above everything the score decided, and a recalculation will not move it. "Clear my order" gives the computed ranking back.'

The '22 opportunities' section shows '0 on the shortlist' and buttons for 'Find unadvertised roles', 'Compare selected', 'Generate applications for the selected', and 'Delete selected'. The first opportunity is 'AI Engineer' at 'BLUECRUX', marked as 'Advertised vacancy'. It includes details about the employer, location, and salary, along with a 'Dream fit' score of 59 and a 'Why this rank?' section showing '9 met', '14 partial', and '3 violated' criteria.

opportunities · the ranked list, filtered and sortable

Why a role ranks where it does

The detail screen turns the score into an argument: seven sub-scores, the dream-job criteria, the compensation estimate and its sources, and the route to a person.

- Sub-scores: profile fit, dream-job fit, directive fit, company, compensation, plausibility, reachability.
- Dream-job fit shows met, partly met and violated criteria, not a single number.
- A speculative opening explains why the system thinks the role exists (FR-263).
- 'Prepare an application for this role' is the bridge into Apply.

The screenshot displays the 'Opportunities' section of the Dream Job application. The main content area shows a job listing for 'Documentation engineer M/V/X' by NOVA Engineering. Below the job title, there are several sections: 'Why it ranks here', 'Sub-scores', 'What each component actually looked at', and 'Dream-job fit'.

Why it ranks here

The employer is not disclosed; this vacancy was posted by NOVA Engineering on behalf of an employer it does not name, so the company dimensions were not assessed. Computed without a narrative model. Strongest components: plausibility 100%; dream fit 50%. Weakest: profile fit 6%. See the sub-scores for the comparisons behind each number.

Sub-scores

Sub-score	Score
Profile fit	6
Dream-job fit	50
Directive fit	37
Plausibility	100
Reachability	10

What each component actually looked at

Profile fit **deterministic** 5.8

- 0 of 3 required skills are evidenced in the profile
- domain overlap on documentation, legal

Directive fit **deterministic** 37.2

- location not precise enough to place
- work arrangement not stated
- permanent contract matches; 100% FTE is within the directive
- posted by a staffing agency; you prefer direct employers; size, stage, ownership and trajectory cannot be checked; the employer is not named, and these describe the agency

Company **not assessed**

- employer not named - cannot assess; this vacancy is posted by an agency
- not scored against the intermediary; its accounts say nothing about the employer's ability to pay
- the posting says: Onze klant is een internationale speler gespecialiseerd in hoogwaardige procesinstallaties en oplossingen voor onder meer de farmaceutische, chemische en voedingsindustrie.

Dream-job fit

Violated: 0 Partly met: 3 Met: 0 Not recorded either way: 16

Partly met

- Target role: Senior Full-Stack / Web Application Engineer** - Strong
partly overlaps the target role 'Senior Full-Stack / Web Application Engineer'
- Target role: AI / Data Engineer (product engineering side)** - Strong
partly overlaps the target role 'AI / Data Engineer (product engineering side)'
- Target role: Technical Lead / Principal Engineer with design-adjacent strategy remit** - Strong
partly overlaps the target role 'Technical Lead / Principal Engineer with design-adjacent strategy remit'

Not recorded either way

- Target role: Hands-on Software Architect** - Must
no textual overlap with 'Hands-on Software Architect'
- Role family: hands-on software engineering & architecture** - Strong
cannot tell whether this is 'hands-on software engineering & architecture'
- Role family: web application development** - Strong
cannot tell whether this is 'web application development'
- Role family: AI & data engineering** - Strong
cannot tell whether this is 'AI & data engineering'

Writing production code personally, as a normal part of the job - Must
the description does not mention this activity

Shaping and owning architecture decisions - Must
the description does not mention this activity

Applying object-oriented design to software being built - Must
the description does not mention this activity

Designing and building well-structured web applications - Must

opportunity detail · the full explanation

Companies and the knowledge base

The shared knowledge base: company profiles, five-year financial analyses, competitors and hiring signals — searchable without picking a campaign.

- Company research is shared across job seekers and reused while fresh, so a second search costs nothing.
- Freshness is shown before the figures: a stale row describes the company as it was.
- Suppression and watch controls, plus an employer-kind verdict for who is hiring.
- The inventory today holds 5,941 companies and 57,293 vacancies.

Dream Job
Stephane van der Aa

DO THIS

- Start here
- 1 - Your profile
- 2 - Opportunities
- 3 - Apply
- Follow up

ADVANCED

PROFILE & SEARCH

- Composite profile
- Dream job
- Directives
- Campaigns
- Browser session

RESEARCH & APPLY

- Companies**
- Dream-job gap
- Contacts
- Apply browser
- Networking

JOURNEY & FOLLOW UP

- Where I am
- Responses
- What works

SYSTEM

- Monitoring
- Mail setup
- Administration

Companies

Companies 5,941 Jobs 57,293 Contacts 3,959 stephane@stepvda.com ? Help Sign out

Profile Plan **Discover** Apply Follow up

The shared knowledge base of company profiles, five-year financial analyses, competitors and hiring signals — searchable on its own, without picking a campaign.

Search the knowledge base

Name, summary, product, market... BE Any size Any stage Any trajectory 62.01 or NACE Any source Search

5941 companies on record - 3279 reached through their own ATS board - mostly from board.eures (1359), ats.personio (1324), board.arbeltnow (884)


5941 companies shared across every campaign 0 on your watchlist

COMPANY	COUNTRY	SECTOR	SIZE	STAGE	TRAJECTORY	FRESHNESS	Watch	Suppress
lonQ Inc. greenhouse.io · SIC 7373 Services-Computer Integrated Systems Design...	US		-	listed	growing	fresh today · 13 Sept 2026	Watch	Suppress
DIGNIFY dignify.be · Legal form: Private limited company; Register status: active; Registered since May 18, 201...	BE		51-200	established	-	fresh today · 13 Sept 2026	Watch	Suppress
Sumble Inc workable.com · Sumble Inc builds a knowledge graph from web data, with an initial focus on data for go-to...	US	Web data and go-to-market intelligence software	11-50	startup	-	fresh today · 13 Sept 2026	Watch	Suppress
Heartbeat AI heartbeat.ai · Heartbeat AI GmbH describes itself as the maker of an emotion-recognition and text-analyti...	DE	Computer programming activities, Web portals	-	startup	growing	fresh today · 13 Sept 2026	Watch	Suppress
Omega omega.de · Omega is a German deep-tech company developing non-invasive, purely optical in-ovo sex de...	GERMANY	Agtech / animal welfare technology (in-ovo sexing for poultry), Other research and experimental development on natural sciences and engineering	-	startup	growing	fresh today · 13 Sept 2026	Watch	Suppress
Cosuno cosuno.de · Cosuno is a Berlin-based software company that provides a software-as-a-service platform f...	GERMANY	Computer programming activities, Construction technology (contech) software	51-200	scaleup	growing	fresh today · 13 Sept 2026	Watch	Suppress
OST BAU ostbau.de · OST BAU is a German construction company that describes itself as a medium-sized building...	GERMANY	Construction of residential and non-residential buildings, Construction of roads and motorways	501-1000	established	growing	fresh today · 13 Sept 2026	Watch	Suppress
Michelscom GmbH michelscom.de · Michelscom GmbH is a German	GERMANY	Tätigkeiten von Arbeitsvermittlungsagenturen, Personalmarketing und Recruiting-Dienstleistungen	11-50	established	growing	fresh today · 13 Sept 2026	Watch	Suppress


companies · the shared inventory

Inside a company


Five years of filed figures become two advisory scores that cite their inputs; market timing, peers and values are separate tabs.




company · Financials — five-year analysis and two scores



company · Market and timing — signals and peers to adopt



company · Values and culture — what the company says about itself



company · the header: trajectory, freshness, watch and suppress

Dream-job intelligence — the gaps

How far the market and your profile are from the job you described, and what would close the distance. Gap analysis and paths work without the language model.

- Gaps are computed against the structured dream-job model, with a threshold the job seeker controls.
- Each gap says what closes it and where it was decisive on the ranked list.
- The reading is stored; recompute is explicit, nothing recalculates behind your back.
- The numbers are deterministic; a model only sharpens the wording (CR-405).

The screenshot displays the 'Dream-job intelligence' interface. The top navigation bar includes links for Companies (5,942), Jobs (57,296), Contacts (3,959), and a user profile (stephane@stepvda.com). The main content area is titled 'Dream-job intelligence' and features a progress bar with steps: Profile, Plan, Discover (active), Apply, and Follow up.

Below the progress bar, a section titled 'How far the market and your profile are from the job you described, and what would close the distance. Five readings of the same question, on five tabs.' shows tabs for 'What works' and 'Dream-job intelligence' (selected). A yellow box contains a disclaimer: 'A reading of the market, not a verdict on you. Every gap, route and mismatch below is derived from your own profile and from the postings and company pages this campaign actually collected. Nothing here is a statement about you that the material does not support, no score decides anything on its own, and nothing changes your ranked list or your profile unless you ask it to. These figures rank and suggest; they never decide. Whatever they point at, the choice to act — or not — stays yours.'


The interface shows a campaign 'Search 11 Sept - completed' with tabs for Gaps (12), Stepping stones (2), Fit across the market, Values conflicts, and LinkedIn profile. The 'Gaps' tab is active, displaying 'What stands between your profile and the dream job ?' with a 'Recompute' button. The results show 12 gaps between the profile and the dream job, across 3 of the six dimensions. The largest gap is 'dream-job activity not evidenced: shaping the architecture of the systems being built, which costs points in 8 of the 72 collected opportunities. 5 of them are blocking. 92 opportunity links were computed from the stored sub-scores. Computed 13 Sept 2026 deterministic 12 gaps - 12 with linked opportunities'.

Below this, a section titled 'Dream-job activity not evidenced: Shaping the architecture of the systems being built' (marked as blocking) shows the dream-job statement and a 'WHAT CLOSES IT' section. It suggests taking a piece of work that is genuinely about shaping the architecture of the systems being built and see it through to a result worth writing down. It also shows 'Where it was decisive' with two examples: 'Full Stack Engineer, AgentControl' (losing 100 profile-fit points) and 'DeepTech Senior Full-Stack Engineer (100 % remote) (m/f/d)' (losing 50 profile-fit points).


intelligence · Gaps and what closes them

Stepping stones, fit and values


When nothing clears the dream-job threshold, the system proposes two or three sequences of roles that lead there — and shows how much of the market clears it at all.




intelligence · Stepping stones — routes to the dream job



intelligence · Fit across the market — the distribution and threshold



intelligence · Values conflicts — where a company contradicts you




intelligence · LinkedIn profile — suggestions you copy by hand

Hiring contacts

The people behind the opportunities, found by walking a ladder of sources and validated rather than guessed.


- A company showing 'unreachable' is a real finding: nothing survived the ladder.
- Contacts are shared unless collected by browser automation, then restricted (NFR-302/303).
- An objection blocks the address permanently for everyone on the installation.
- The knowledge base holds 3,959 contacts today.




contacts · the ranked ladder for a company

Contact discovery at scale


Browse every contact, find missing addresses, walk warm introduction routes, import your network and honour objections and retention.




contacts · Browse all — filters and the missing-email panel



contacts · Introduction routes — warm paths into a company



contacts · My network — import your own connections




contacts · Objections and retention — the privacy sweep

Applications and the Apply browser

The generated CV, briefing, motivation document and email for each shortlisted role — all for the job seeker to read, edit and approve.


- A consistency check catches invented employers, dates and titles (FR-322).
- A leak scan refuses briefing and motivation as attachments (NFR-206).
- Approving authorises dispatch but sends nothing; messages go from the seeker's own mailbox.
- The job seeker is the last line of defence for anything technically true and still wrong for the audience.




apply browser · the email, with every review tab

The four artefacts, and the checks before dispatch


CV, briefing, motivation and email are generated together; the briefing and motivation never leave the machine, and the checks gate the send.




apply · CV — grounded in the profile, nothing untraceable



apply · Briefing — for the job seeker, never attached



apply · Motivation — the longer statement, never attached




apply · Checks — consistency and leak scan before dispatch

Networking and export

Warm introduction routes, an event radar with calendar export, and a self-contained campaign package for a coach — built from one structure, redacted like the

- Introduction routes are ranked by strength and relevance (FR-461).
- Events come with a match meter, calendar links and an .ics download (FR-462).
- The export is a PDF bundle plus machine-readable JSON, and it refuses outright if it could reach another seeker's data (FR-463).
- Discretion mode is honoured on every response (FR-385).




networking · Event radar and the export builder

Application pipeline

Every authorised application and its stage, from sent to reply, interview and outcome, on one board.

- Open a card for the classified reply, the drafted answer and the stage dates.
- Approving a draft queues it for the mail screen; it still sends nothing.
- Outcomes feed back into generation defaults and scoring weights.
- The negotiation brief and a mock interview sit on the same card.




pipeline · the five-stage board

Responses received

Every answer to an application, however it arrived: detected automatically from Gmail, or recorded by hand for a call, a portal or a LinkedIn message.

- Recording rejections matters as much as good news; both feed the pattern analysis.
- An application counts as answered or as silence, and nothing in between until 21 days pass.
- Recording a response moves the card on the pipeline board automatically.
- A wrong outcome is correctable, because it distorts every rate it is counted in.




responses · record, correct and review

What works, and where to redirect

Which kinds of job and company actually answer, computed from recorded outcomes and always shown with the sample size behind them.

- Rates are drawn with a plausible range, so a thin segment cannot look decisive.
- Advice needs roughly six resolved applications before it will say anything.
- Applying a proposal creates a new directive version — it never overwrites (FR-285).
- The data does not get to overrule the dream job; it only shows what the mismatch costs.




insights · outcome rates with Wilson intervals

Monitoring

The operational surface: campaign health, token spend, extraction rates and the audit trail, plus the scheduled re-checks of watched companies.

- Extraction-rate monitoring surfaces adapter breakage (NFR-403).
- Every LLM call, egress request and score is audited with its provenance.
- Watched companies are rechecked on a schedule, and what changed becomes one weekly digest (FR-401/403).
- Notifications, timing windows and the scheduler are all on this screen.




monitoring · watchlist, digests and the scheduler

Mail setup

Two ways to send from the seeker's own identity — Gmail OAuth or a Resend relay — behind one dispatcher that holds every guard rail.

- Gmail and Resend both satisfy 'send from your own mailbox' (FR-325).
- Sending rules, windows and a daily cap protect sender reputation (NFR-702).
- The dispatch log records everything sent, down to the message id; replies are detected.
- Follow-ups become due on a schedule and are drafted, then approved, never auto-sent (FR-327).




mail setup · mailboxes, rules, log and follow-ups

Administration

The installation-level controls: the source catalogue and its terms acknowledgements, the server logs, model routing and the data-rights tools.


- Nine tabs: Users, Models, Sources, Activity, Continuous, Audit, Logs, Data, Employer kind.
- Sources that require acknowledgement stay disabled until an administrator confirms their terms (IR-101).
- Model routing chooses the provider per task and can point at a local endpoint for private data (FR-362, NFR-306).
- Two actions cannot be undone: the redaction sweep and erasing a job seeker; erasing leaves the shared market data standing (FR-108).




administration · the source catalogue and its terms

Administration, the other eight tabs


Activity and Continuous show what the installation has collected; Models controls spend and routing; Data holds the irreversible actions.




admin · Models — provider, routing, budget and usage



admin · Activity — corpus, traffic and cost counters



admin · Data — retention, redaction and erasure




admin · Users — accounts, admin role and password resets

Help, everywhere

Every screen explains itself: inline '?' tips drawn from a 142-term glossary, a per-screen drawer with purpose, steps and cautions, and first-run guidance that works

- Press ? anywhere to open the drawer for the screen you are on.
- The drawer carries purpose, how to work through it, worth knowing, and a caution where reliance is risky.
- Inline tips resolve from the same glossary, so a term means the same thing everywhere.
- Help is static and offline: it works before a profile exists.



help drawer · opened with the ? key on any screen


PART 2

Technical architecture

The layered system, the five chokepoints that make its requirements enforceable, and the path one campaign takes through the pipeline.

The shape of the system

A React single-page application over a FastAPI API, a pipeline of stages, and five services — LLMClient, EgressClient, SourceAdapter, JobRunner and crypto — over a repository layer and SQLite.




422 routes over 88 tables · 227 Python modules, 140 JS modules · 2,061 tests.

The five chokepoints

The architecture funnels five concerns through exactly one module each, which is what makes the requirements below them enforceable rather than aspirational.

- Every SQL statement — `db/repositories/*`, `db/connection.py`: portability and isolation (CR-408, FR-344).
- Every outbound HTTP request — `egress/client.py`: robots, pacing and raw capture (FR-182/183).
- Every LLM call — `llm/client.py`: budget, injection defence, redaction, routing, audit.
- Every external source — `adapters/base.py`: one file per source, terms status (NFR-601, IR-101).
- Every long-running task — `jobs/runner.py`: pause, cancel, resume (FR-185, NFR-401).




Requirement identifiers as cited in the code that enforces them.

The path of one campaign

From directives to redirection advice. The middle of the flow reads and writes the shared knowledge base; replies stay private to one job seeker.

- Plan: directives and the composite profile become a priced, per-source plan with a human gate before collection (FR-161..166).
- Reuse: the planner asks the knowledge base what is fresh enough and collects only the rest (FR-342/343).
- Collect: the plan runs through the adapters and EgressClient into raw documents and knowledge-base rows (FR-181..186).
- Analyse and rank: companies, filings, opportunities, compensation and seven sub-scores produce the advisory ranking (FR-221..285).
- Apply and learn: contacts, documents, dispatch, replies and redirection advice that becomes the next directives (FR-301..285).



One campaign, stage by stage

Every stage persists its own artefacts and can be re-run on its own (NFR-603). The user-facing pipeline is the five phases: Profile, Plan, Discover, Apply, Follow up.


Stage	What it produces	Requirements
Plan	Campaign plan: sources, queries, cost	FR-161..166
Reuse	What the knowledge base already has, fresh enough	FR-342/343
Collect	Raw documents through the adapters and egress	FR-181..185
Profile	Shared company profiles	FR-221..225
Financials	Five-year analysis, ability to pay, investment capacity	FR-241..246
Enrich	Hiring signals, employer kind, domains	FR-201..207
Synthesise	Opportunities, including speculative openings	FR-261/262
Price	Compensation estimate against the stated minimum	FR-264
Score	Seven sub-scores and the ranked list	FR-281..285
Contacts	Hiring contacts, validated and ranked	FR-302..306
Generate	CV, briefing, motivation and email	FR-321..324
Dispatch	Send, approved and from your own mailbox	FR-325..327
Learn	Replies, outcomes and redirection advice	FR-285, FR-425

The adapter contract

One contract — plan, fetch, parse, normalise — and every source is one subclass and one decorator (NFR-601).


- plan: directives, composite profile and caps become this source's own query form, priced and paced.
- fetch: every request goes through EgressClient — robots, rate limit, cache and raw capture.
- parse: JSON-LD first, then deterministic selectors, then LLM extraction, in that order.
- normalise: knowledge-base columns carrying a confidence and a provenance record (NFR-402).

One contract



NFR-601: a new source is one subclass and one decorator.

29 adapters, by type



4 of them ship disabled until an administrator acknowledges their terms (IR-101).

The path of one LLM call

Every call in the system takes this path, because llm/client.py is the only module that speaks to a model.

- Redact, route, fence: do-not-disclose fields stripped, a model chosen per task, scraped text quarantined (CR-410, NFR-205).
- Check the budget pre-flight, then shed optional work rather than stop mid-campaign (NFR-104).
- Validate the response before any of it is used; injection defence is at the boundary.
- Debit and audit: tokens, cost, latency, model and prompt version recorded (FR-364).

One LLM call, end to end

Every call in the system takes this path, because llm/client.py is the only module that speaks to a model.

```
llm.complete_json("extract.vacancy", system=..., user=...,  
untrusted={"page": html}, entity_type="vacancy")
```

the task id is what routing, cost attribution and the audit row all key off


1	Redact	do-not-disclose fields and special-category data stripped before anything leaves	CR-410 FR-106 · FR-127
2	Route	administrator setting, then a per-task model, then the cheap / strong split	FR-362
3	Route · privacy	composite profile, tailored CV and motivation to a local endpoint when one is configured	NFR-306
4	Fence	scraped text goes in labelled untrusted blocks, never into the instruction	NFR-205
5	Check the budget	pre-flight estimate; shed optional work rather than stop mid-campaign	NFR-104
6	Call the provider	OpenAI-compatible POST to DeepSeek, or to the local model	CR-409
7	Validate	the response is parsed and checked before any of it is used	NFR-205
8	Debit and audit	tokens, cost, latency, model and prompt version written to llm_call	NFR-104 FR-364

24 versioned prompt templates; the version travels into llm_call, so a bad output can be tied to the revision that produced it.

One outbound request, end to end

No other module opens an HTTP connection. Cache, robots.txt, pacing, fetch, capture, parse — every response kept under its content hash (FR-183).

- A cache hit never reaches the network.
- robots.txt is parsed per domain; a disallow raises rather than being ignored.
- A 429 or 503 backs off exponentially under one retry policy (IR-102).
- Raw bodies are stored on disk so an extractor can be re-run on the page it originally saw.



FR-183 · every response is kept under its content hash, so an extractor can be re-run on the page it originally saw.


One database, three scopes

49 private tables carry a `job_seeker_id`; 38 shared tables carry none; the migration ledger is the third. FR-344 is the invariant: no shared row links back to a job seeker.

- Private: profile versions, campaigns, opportunities, applications, replies.
- Shared: companies, vacancies, financial years, signals, competitors, events.
- Restricted: contact straddles the boundary — browser-collected rows are private to the campaign that made them (NFR-303).
- Two FTS5 virtual tables index companies and vacancies; a `usable_contact` view encodes the objection rule at the database.
- Erasing a job seeker deletes the private rows and leaves the shared market data standing (FR-108).

One database, three scopes

88 tables: 49 carry a `job_seeker_id`, 38 carry none, 1 records the migrations.




FR-344 is the invariant: no shared row carries a job-seeker id, so erasing a job seeker leaves the market data standing.

Every long task is resumable

A job runner is the single home for long tasks: jobs are pauseable, cancellable and restart-safe, and they never run on the request loop.


- Jobs run on their own threads with per-unit checkpoints (FR-185, NFR-401/502).
- A paused job resumes where it stopped, not from the beginning.
- All job threads use the bulk write lane, so requests never queue behind a campaign (NFR-102).
- Job outcomes, progress and errors are persisted and surfaced on screen.



In-process async on a single node: no broker, one writer. A crash loses at most the page in flight.

The codebase, measured


The pipeline is the centre of gravity; the interface and the tests are the next two largest bodies of work. Numbers are measured, not estimated.



the pipeline alone is 40% of the backend
16 packages · 422 API routes · 88 tables

Where the lines live

226,209 lines in all



tests are 27% of everything written

2055 tests passing


2,061 tests collected across 114 files; tests are 34% of the backend's lines.

What the interface costs

Route-level code splitting keeps the first paint small; the heaviest screen loads only when it is opened.

- Before splitting the bundle was 858 kB; the shared entry is now 312 kB, 105 kB over the wire.
- 33 lazy chunks: a screen's code downloads when the route is opened (NFR-101).
- Every phase hue meets WCAG AA for small text, on white and on its own tint, in both themes.
- The screenshots in Part 1 are the running application, not mock-ups.

Before — one bundle, everything up front




After — a shared shell, routes on demand



The administration area and the PDF previews are no longer in the first payload.
Nothing is downloaded to reach the profile page but the profile page (NFR-101).

Requirement coverage

156 of 157 requirements in the specification are cited in the implementation — 99.4%; the one not cited is the DPIA (NFR-304), a document rather than code.



The two not cited in code:

NFR-304 (Must) · A DPIA shall be produced before production use — a document, delivered as DPIA.md


Must 103/104, Should 44/44, Could 9/9. The line is not cited is NFR-304.

Seven sub-scores, one explained rank

Every rank is a weighted mean over seven named sub-scores; the weights are visible, editable and learned from outcomes, and the score orders a list rather than making a decision.

- Profile fit 0.22, dream-job fit 0.22, directive fit 0.18, company 0.14, compensation 0.10, plausibility 0.08, reachability 0.06.
- Deterministic pre-rank of the whole corpus, then LLM scoring of the top 500 by default.
- The dream-job meter uses must/strong/nice importance and met/partial/violated status.
- Scores are advisory (NFR-305); manual order and pins win over the computed rank.

Seven weighted sub-scores




A role can fit a CV perfectly and still not be what the person wants, so the meter is reported beside the score and never inside it.

Learning it refuses to overstate

Outcome rates are shown with a Wilson interval and the sample size, and the system will not advise below six resolved applications.

- Every rate carries its interval and its n , so a thin segment cannot look decisive.
- Advice needs roughly six resolved applications before it speaks (FDD 9.3).
- Applying a proposal writes a new directive-set version; nothing is auto-applied (NFR-305).
- Rejections are data too: they feed the same segmentation as interviews.




Every point is the same observed rate: 50%. What changes is the 95% interval around it:
81 percentage points wide at $n = 2$, 21 at $n = 80$.

Two advisory scores from filed accounts

Five years of filed figures are reduced to ability to pay and investment capacity — weighted, coverage-shaded, and never claiming more than the evidence supports.

- Ability to pay: cost per FTE 0.35, EBIT margin 0.25, equity/assets 0.20, current ratio 0.20.
- Investment capacity: revenue CAGR 0.22, cash runway 0.22, headcount CAGR 0.18, EBITDA margin 0.18, debt/equity 0.12, capex/revenue 0.08.
- Components that cannot be computed are not invented; the score is shaded toward 0.5 for the missing coverage.
- Both scores are advisory and cite the figures they rest on (NFR-305, RK-06).



A profile built from secondary signals rather than filings can never score above 55 (FR-245).

PART 3

Data sources

Every adapter registered in the codebase — 29 today — what each one covers, how it is reached, and whether it needs consent.

The source catalogue at a glance

29 adapters register themselves at import time. 4 require an explicit terms acknowledgement before they can be enabled (IR-101).

Group	Count	Consent	Access
Applicant tracking systems	9	1 require acknowledgement	API, API, API...
Job boards and public employment services	9	2 require acknowledgement	HTTP, API, API...
Company registries	5	—	API, HTTP, API...
Directories	1	—	API
News and signals	1	—	HTTP
Conferences and events	2	1 require acknowledgement	HTTP, HTTP
Company websites	1	—	HTTP
Compensation benchmarks	1	—	API

Nothing is collected from a source whose terms are not acknowledged; the planner selects from this catalogue using the coverage each adapter declares (FR-161/164).

Applicant tracking systems

Structured job feeds hosted by the employer's own recruiting system. Machine-readable, so collection costs no tokens.

Source	Type	Access	Terms	Acknowledgement	Enabled
Ashby	Applicant tracking systems	API	permitted	—	yes
Greenhouse	Applicant tracking systems	API	permitted	—	yes
Lever	Applicant tracking systems	API	permitted	—	yes
Personio	Applicant tracking systems	API	permitted	—	yes
Recruitee	Applicant tracking systems	API	permitted	—	yes
SmartRecruiters	Applicant tracking systems	API	restricted	required	yes
Teamtailor	Applicant tracking systems	API	permitted	—	yes
Workable	Applicant tracking systems	API	permitted	—	yes
Workday	Applicant tracking systems	API	permitted	—	yes

Job boards and public employment services

Advertised vacancies across Belgium, the Netherlands, the EU and globally. Terms status decides whether a source is enabled by default.

Source	Type	Access	Terms	Acknowledgement	Enabled
Actiris (Brussels public employment service)	Job boards and public employment services	HTTP	permitted	—	yes
Arbeidnow	Job boards and public employment services	API	permitted	—	yes
EURES (European Job Mobility Portal)	Job boards and public employment services	API	permitted	—	yes
Generic HTML job board	Job boards and public employment services	HTTP	permitted	—	catalogue only
Indeed	Job boards and public employment services	HTTP	prohibited	required	yes
Jobat.be	Job boards and public employment services	HTTP	restricted	—	yes
StepStone	Job boards and public employment services	HTTP	restricted	required	yes
VDAB (Vind een job)	Job boards and public employment services	HTTP	restricted	—	catalogue only
Welcome to the Jungle	Job boards and public employment services	API	restricted	—	catalogue only

Registries, directories, news, events, websites, compensation

The non-vacancy sources: statutory filings that price a company, news and events that signal hiring, the company's own site, and the pay benchmark.

Source	Type	Access	Terms	Acknowledgement	Enabled
Companies House (UK)	Company registries	API	permitted	—	yes
KBO/BCE (Belgian enterprise register)	Company registries	HTTP	permitted	—	yes
KvK Handelsregister (NL)	Company registries	API	permitted	—	yes
NBB Central Balance Sheet Office	Company registries	API	permitted	—	yes
SEC EDGAR (US)	Company registries	API	permitted	—	yes
OpenCorporates	Directories	API	permitted	—	yes
Company newsroom / RSS	News and signals	HTTP	permitted	—	yes
Conference sites and community calendars	Conferences and events	HTTP	permitted	—	yes
Meetup and Eventbrite (public listings)	Conferences and events	HTTP	restricted	required	yes
Company website	Company websites	HTTP	permitted	—	yes
Eurostat Structure of Earnings Survey	Compensation benchmarks	API	permitted	—	yes

What the catalogue implies

The catalogue is not a fixed list: the planner selects from it using the coverage metadata each adapter declares (FR-161/164).

- Coverage is declared per adapter, so a Belgian search runs the registries and boards that cover Belgium and skips the rest.
- Terms status is first-class: SmartRecruiters, Indeed, StepStone and Meetup/Eventbrite require acknowledgement before they are enabled (IR-101).
- Adapters that return machine-readable payloads set `llm_fallback = false`, so the structured forms are read deterministically at zero token cost.
- Every source is reached through EgressClient, so robots.txt, pacing and raw capture apply no matter which adapter asked.
- The shared knowledge base is reused while fresh, so a second search over the same companies costs little or nothing (FR-342).

The system in five sentences

Profile in, ranked and explained opportunities out, every message reviewed and sent by the job seeker.

- Profile in: a LinkedIn export, a CV and a description of the job actually wanted, merged into one composite profile with every conflict surfaced.
- Plan and collect: structured directives become a per-source plan, reviewed and priced before anything runs, then executed through 29 registered adapters behind one contract.
- Analyse and rank: company profiles, five years of filings and seven sub-scores produce an explained, advisory ranking of advertised and speculative opportunities.
- Apply with consent at every gate: generated documents are checked, approved by the job seeker, and sent from the seeker's own mailbox — nothing leaves on its own.
- Follow up and learn: replies, outcomes and negotiation are captured; redirection advice becomes the next directive version, which the seeker accepts or dismisses.

Dream Job • v2

Every screen, every diagram and every count in this deck was taken from the running system in September 2026.

227 Python modules • 140 JS modules • 422 API routes • 88 tables • 29 adapters • 2,055 tests passing